

## L01: Introduction to Planning Algorithms

Planning Algorithms in AI and Robotics

Prof. Gonzalo Ferrer

Skoltech, 29 October 2021



# Presentation: Who are we?

**Instructor:**

Prof. Ferrer (name Gonzalo)

(g.ferrer@skoltech.ru)

**Teaching Assistants:**

Aleksandr Gamaiunov

(aleksandr.gamaiunov@skoltech.ru)

Timur Akhtyamov

(timur.akhtyamov@skoltech.ru)

**Mobile Robotics Lab:** Path planning, Robot Navigation in dynamic environments, Pedestrian Motion prediction, Sensor fusion of Lidar, camera, IMU, etc., SLAM, Localization, Mapping, etc

<https://sites.skoltech.ru/mobilerobotics/>



# What is planning? Robotics

**Robot** converts high-level specification of tasks into low-level descriptions of how to move.

Mostly finding a plan is known as **motion planning** or **path planning**.


Example: The *piano movers problem*: how to move one piano from one room to the next room.



# What is planning? Robotics

Other examples, moving an object from a starting configuration to goal configuration while avoiding obstacles.


On this first view we only consider **feasibility**, although later we will consider optimality and uncertainty.



# What is planning? Robotics

Now, image the “piano” as the green rod in this 3D maze.

Dynamics will not be covered on this course, although it is a natural extension to planning (trajectory planning).

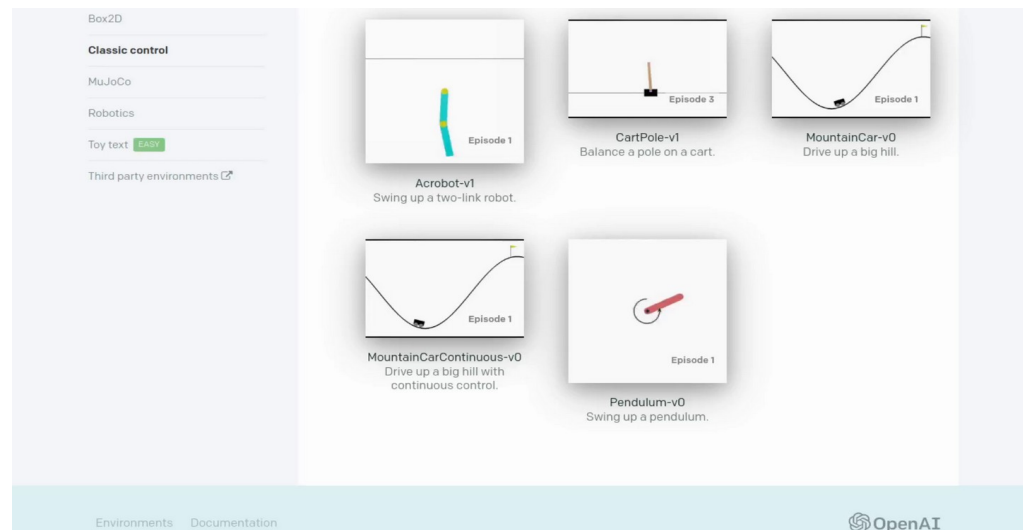


# What is planning? Control Theory

Control Theory typically considers physical systems and differential constraints.

The **controller** designs **feedback policies** or **control laws**.

Controls usually focuses on optimality and stability.



# What is planning? Artificial Intelligence

**Agents or decision-makers** in AI do planning or **Problem Solving**.

Usually discrete spaces, leading to combinatorial solutions.

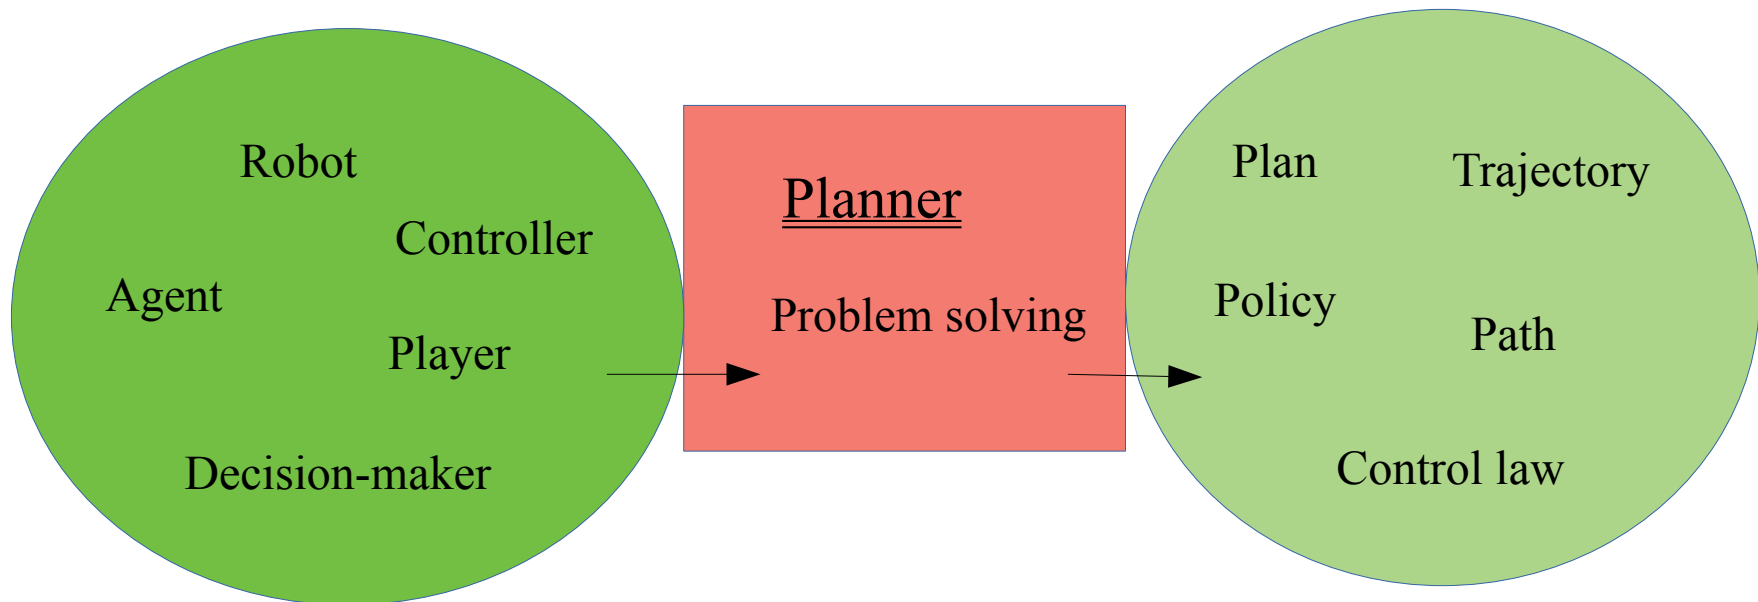
Example: Solving the Rubik's cube.



# Planning to plan

Planning means different things in different disciplines: robotics, artificial intelligence and control theory.

In this course we will present a unified view:



This view is from LaValle's book, Ch.1



# The ingredients of planning

- **State:** All possible situations that could arise.
- **Action:** Change the states. Also known as *inputs*, *controls*, *operations*, etc.
- **Initial and Goal states:** A planning problem involves these 2 states.
- **Criterion:** The desired outcome of a plan:
  - Feasibility: arrival at a goal state, regardless of efficiency.
  - Optimality: Find a feasible plan and optimize some function.
- **Plan:** A specific strategy of behavior. It could be a sequence, a policy, etc.



# The ingredients of planning

**Example:** the Piano movers problem.

State: 3D poses	Actions: 3D (relative) poses
Initial state: the corridor	Goal: my office
Planner: RRT (L04 algorithm)	Plan: sequence of actions


**Example:** How to arrange the furniture in your new house

State: Position of all objects	Actions: Move objects
Initial state: Empty space	Goal: Ikea picture
Planner: Yourself	Plan: sequence of actions.




# Taxonomy of Planning

There are different ways to use the plan calculated by the planner




Open loop execution



Feedback



Refinement



Hierarchical



# Course Goals

Mastering a set of core planning algorithms, covering a wide variety of planning problems, such as discrete planning, continuous planning, decision-making, planning under uncertainty, learning-based, etc.

For any give task, select the right planning algorithm.

# Prerequisites

Basic programming skills in Python

Probability, introductory course.

**Next steps** -> Perception in robotics in T3

<https://github.com/MobileRoboticsSkoltech/Perception-in-Robotics-course-T3-2021-Skoltech>



# Learning Outcomes

- The student will acquire **theoretical knowledge** and a rich set of **practical skills** to design their own planning and decision-making algorithms, applied to any kind of problem related to AI, Robotics, etc.
- The student will be able to **analyze** problems under the perspective of planning, and provide a more diverse set of tools for problem solving.
- The student will get **experience** on different planning techniques such as discrete planning, continuous planning, decision-making, planning under uncertainty, learning-based. We will prepare materials, seminars and problem sets that will serve as a first step into each of these particular topics.



# Course Structure

10 lectures (Online)	Monday 16:00-18:00 Friday 16:00-18:00
70% Problem Sets	PS1: Discrete planning PS2: Sampling-based planning PS3: Value Iteration PS4: Decision-Making
30% Final Group Project	(more in canvas and later in class)



## Course summary:

Date	Details	Due
Fri, 29 Oct 2021	 L1: Introduction	0:00
Mon, 1 Nov 2021	 L2: Discrete Planning	0:00
Mon, 8 Nov 2021	 L3: Configuration Space	0:00
Fri, 12 Nov 2021	 Seminar 1: Distances	0:00
Sun, 14 Nov 2021	 PS1: Discrete Planning	due by 23:59
Mon, 15 Nov 2021	 L4: Sampling-Based Planning	0:00
Fri, 19 Nov 2021	 Seminar 2: Sampling	0:00
Sun, 21 Nov 2021	 PS2: Sampling-based planning	due by 23:59
Mon, 22 Nov 2021	 L5: Discrete Optimal Planning	0:00
Fri, 26 Nov 2021	 L6: Continuous Optimal Planning	0:00
Sun, 28 Nov 2021	 PS3: Value Iteration	due by 23:59
Mon, 29 Nov 2021	 L7: Decision Making and Games	0:00
Fri, 3 Dec 2021	 L8: Markov Decision Process	0:00
Sun, 5 Dec 2021	 PS4: Decision making	due by 23:59
Mon, 6 Dec 2021	 Seminar 3: Reinforcement Learning	0:00
Fri, 10 Dec 2021	 Seminar 4: General course discussion	0:00
Mon, 20 Dec 2021	 Final Project Presentation	0:00

# Course Material

- Lecture slides with annotations, uploaded after every class.
- Videos from lectures, upload to canvas/youtube channel.
- Books:
  - S. LaValle, “Planning Algorithms”. Cambridge university press, 2006
  - Sutton and Barto “Reinforcement Learning: an Introduction”, MIT press 2018
- Canvas, selected papers
- Telegram channel





# Class Structure

Lectures will be imparted online via Zoom/offline, blocks of 45 minutes.

Students are encouraged to participate and discuss in class via microphone or chat.

## **Recommendations:**

Participate as much as possible in class!!

Make questions, be engaging, learn more (even in remote mode)



# Problem Sets

- Problem Sets (PS) will be written in Python
- PSs are substantial work and should be worked on during the full allotted time period (each is a 17.5% of your grade).
- There will be a penalty in the max possible grade of -1%/hour for the first 3 hours and then a penalty of -20%/day up to 50% of the grade.

Grade = min (your grade, max possible grade)

- Late submission is based on the last update to canvas.
- Students are encouraged to discuss on PS. Copying code is forbidden. On every PS there will be a section dedicated to Acknowledgments, if any.
- All PS's must be submitted in order to pass the course.



# Course Policies

## **Attendance**

Attendance is highly recommended. In the online format of the class, this will be your chance to ask, discuss and learn as we progress. Later it will be too late.

## **PS Regrade Policy**

If you believe we graded a problem-set or an exam of yours incorrectly, you can submit a regrade request no later than one week after the graded work is originally returned.

## **Academic Integrity**

Reference to Skoltech's policy (see canvas)



# Final Group Project

- Topic (related to the course): Extend a state of the art algorithm, or paper reproduction or implementation on your own settings. We will upload to canvas a list of selected papers for inspiration.
- 3-5 students / group
- Proposal (December 8): a one page document.  
(ask for feedback from the course team)
- Presentation (December 20) 12'+3' questions
- Final project document, on an IEEE template.



# Past projects

## Rapidly-exploring Random Belief Trees for Motion Planning Under Uncertainty




Fig. 4. Preliminary experimental RRBT results with various hyperparameters (size of covariance matrix, num of vertices, initial and goal points).

## Bringing the learning to classic motion planning




Figure 5: Observed area by RRT applied to learned samples with 0.1,0.5,0.7 probability

# Past projects

## Evader-pursuer zero-sum game by MCTS




Fig. 2. Environment. Red circle corresponds to the pursuer position, blue to the evader position, green to the goal and black regions correspond to the walls.

## Comparison of different algorithms for Pacman environment




Fig. 6. Experiments with Q-learning on small environment